01 AI discovery engine

6,551 public posts read. One failure dominates

1 COLLECT
6,551 posts

2 FILTER
1,196 kept

3 READ
15 fields each

4 SORT
84 real cases

5 COMPARE
live tool

How it works: Reddit, Play Store and App Store posts from five countries. A keyword filter drops storage and billing complaints before any AI runs. Gemini then pulls the same 15 facts out of every post — what the photo was, what the person remembered, what they had forgotten, where the search broke, what they did instead — and a second pass separates real memory failures from bug reports.

82% gave the app a fair clue, and it read it wrong

31 of 84 typed one keyword. Only 3 could not say what they wanted

19 remembered words printed in the photo, and still failed

16 had forgotten only one thing: roughly when

Why this is more than reading reviews: each post is tagged to the point in the journey where it broke, so the failures can be counted and sliced by photo type, by what the person remembered, and by what they did next. Every number traces back to a quote. Public posts rarely say what the photo actually was (57% never do) which is why we ran interviews.

DISCOVERY ENGINE
Live and testable. Filter the 84 cases by photo type, by what was remembered, or by where it broke.

Free tiers only: open-source scrapers → Python → Gemini → Streamlit Cloud. No spend.

Hand audit: 20 classifications read and judged by hand, 12 agreed. The misses were nearly all posts that never describe a photo; dropping those leaves 51 cases and moves the 82% to 76%. The finding holds, the precision does not.

02 Metric breakdown

Six steps from memory to match. Step three breaks

“THE ONE” NUMBER Of the times someone hunts a photo over a year old, how often do they end up opening it.

1 **REMEMBER**	They hold pieces: who, where, what was in it.	**90% recall one thing. 68% recall two.**
2 **EXPRESS**	They type something, or give up and scroll.	**74% scroll. Only 48% ever search.**
3 **UNDERSTAND**	The app decides what each word means.	**82% of failures start here.**
4 **FETCH**	It searches for what it thinks you said.	**Works. “august 2025” returned it at once.**
5 **MATCH**	They scan what came back for the photo.	**One query returned it, buried in 64.**
6 **REFINE** ↻	They go back to step 2 with one more clue.	**77% needed several tries. Nothing is kept.**

Nothing came back A dead end, and the most common one.

Some, none right The most useful thing, thrown away.

Found it Opened, in time to matter.

An honest disagreement: External research says step 5 is the bigger prize. Our User Research says step 3. Not yet settled.
Step 6 barely exists today: a near miss is the most useful thing someone can tell us, and the app discards it.

Most people never search. They scroll, and forget the date

74% scroll by date

61% forgot the date they were scrolling to

Only 48% use search feature

52% fail to find a photo monthly or more

THE BIG LIBRARY
Four years, 75,000 items. Fails weekly. Scrolls first, searches last.
"I just wish any agent could find what I need"

THE PAPERWORK KEEPER
Photographs IDs, bills and receipts instead of filing them. 77% do this.
One now keeps documents in Drive, not Photos

THE CASUAL ONE
3,000 photos, six years, no system. Runs into it rarely, so gives up fastest.
"They show up, but not when I want em to"

Paperwork fails for everyone
Three of five could not find a document they had photographed: an ID card, a password screenshot, a bill.

Face search fails as libraries grow
It worked at 5,000 photos and failed at 75,000: an old face no longer matches the person today.

All five work around the app
Take the photo again, file it in Drive, or wait and hope. None of this shows up as a failed search.

How we ran it: we gave people real retrieval tasks on their own libraries and watched, plus a survey of 31. Out of those 31 users we conducted five User interviews.
One idea that came back: four people said juggling several Google accounts never hurt, so we dropped it. The fifth said it did. Back in as an open question.

04 Target segment and size

Who this is for: big libraries and old photos

User Segment that we have taken -> User having:
Account 3 years or older
At least 2,000 items
At least 15% of them paperwork
Tried to find something over a year old, in the last month

WHY NOT EVERYONE
Forgetting depends on how much you have and how long ago it was. Someone with 400 recent photos just scrolls and finds it. The problem starts when the library outgrows memory, which is also when people are paying for storage.
Our interviews show it: face search found a date-blind photo in a 5,000-item library and failed the same task at 75,000.

1.5B
people use Google Photos each month

× 35%
have big, mixed libraries

= 525M
in this group

× 52%
fail to find a photo monthly

= 273M
people failing every month

About 96M of those really matter — 35% of people who fail say it sometimes genuinely counts.

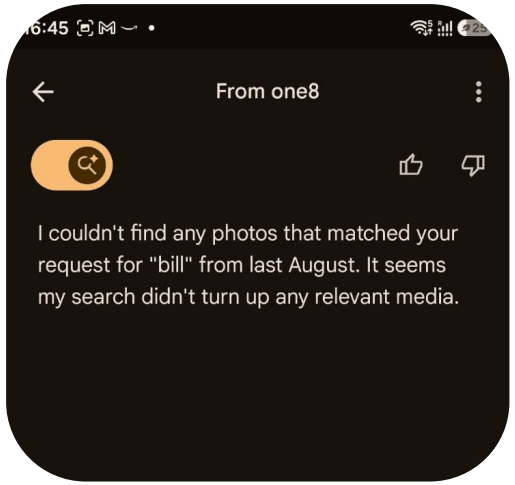
Who we are not building for — under three years or under 2,000 items. They scroll and they find it.

These are estimates, and we say so. The survey is 31 people who volunteered, so we halved the share it suggested. The point is the order of magnitude, not the exact number: even a tenth of this is a lot of people.

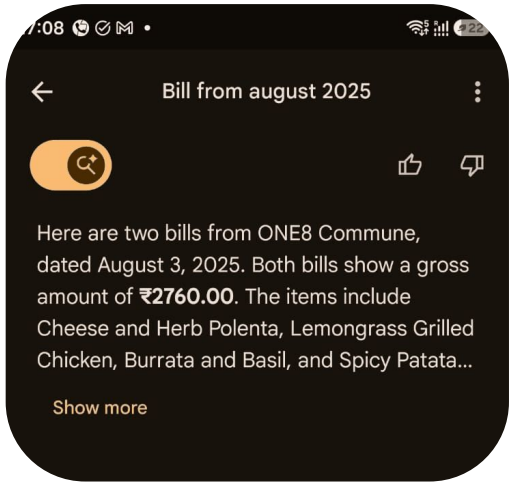
We reproduced it: a wall between your words and its guess



“bill from last august”
FAILED -> got Aug 2026



“... last august at one8”
FAILED -> venue hurt



“bill from august 2025”
FOUND at once. Control.



“... around a year ago”
FOUND, buried in 64

There is a wall between what you typed and what the app decided you meant. You wrote “last August”. It heard “August 2026”, made that a hard filter, and dropped your bill, which was from August 2025. Reading the photo was never the problem: it could list the dishes. Nothing on screen said a guess had been made.

Three ways people say when, and the middle one is the trap: an exact date becomes an exact filter and works · “last August” gets quietly pinned to one year and fails · an openly vague phrase becomes a wide range, and works.
What we built from this: the app shows the words it took from your sentence, lets you change any of them, and asks when one could mean two things.
Live app, with a control, 28 Sep 2026. Screenshots cropped to the reply text.

06 Solution options

Three ways to help people say more before we guess

All three attack step 3, where the test showed it breaks. They differ by what the person can actually give us.

A · THE COMPOSER when you can describe it

Your sentence turns into small editable tags: sister · night · outdoors · wedding? · date unsure. A question mark means a preference, not a filter, and one tap fixes a wrong reading.

Reach 6 · Impact 9 · Confidence 9 · Effort 3

RICE 162

(SELECTED)

B · THE PROMPTER when you do not know what to say

The box stops being empty. It asks for what memory actually keeps, strongest first: who was there, indoors or outdoors, roughly where, what was happening. All optional. None of them a date.

Reach 9 · Impact 7 · Confidence 5 · Effort 4

RICE 79

C · THE ANCHOR when you can only point

“It is from the same trip as this one.” Point at any photo you can find and search around it: same day, same people, just before or after. Words optional.

Reach 6 · Impact 8 · Confidence 8 · Effort 5

RICE 77

REACH how many searches it can help with

IMPACT how much it moves a failing search

CONFIDENCE how strong our evidence is

EFFORT build cost, lower is better

Why A wins, even though B reaches more people. B reaches the 74% who never type anything, and it still loses. It scores 5 on confidence because no user asked for a form — the idea came from us. A wins on evidence: a test where only the date wording changed, plus three unprompted asks from two people for exactly this. B is next, not dropped.

Deliberately left out: fixing the dead end at step 5, which outside research rates first and we rate second. Our testing did not settle that, so it stays open.

07 Choosing one

We Challenged all three. Here is what survived

A · THE COMPOSER

THE BEST ARGUMENT AGAINST

A filter drawer with extra steps, and the research says nobody should build one.

OUR ANSWER

The prompt comes from the person's own words, might not be exactly what a model might need(e.g. a menu. A receipt, not a control panel.)

B · THE PROMPTER

THE BEST ARGUMENT AGAINST

It is a form. Forms feel like work, and nobody asked for one.

OUR ANSWER

True, which is why it scores 5 on confidence. It only shows up after a blank result or a long pause.

C · THE ANCHOR

THE BEST ARGUMENT AGAINST

Google already has “more from this day”. This is browsing renamed.

OUR ANSWER

Browsing from a photo exists. Searching from one does not. P2 did this by hand to fix a date a month out.

Why A holds up on all three things the brief asks about

Does it solve it? “bill from last august” returned nothing; “bill from august 2025” returned it at once. Only the wording changed.
Is it different? Google, Apple and Immich all guess silently. ChatGPT and Gemini ask, but cannot see your photos.
What stops a copy? Anyone could copy the screen in a week. The hard part is knowing whether “last August” is ambiguous for you, which needs your own history.

How big: 273M fail monthly, 82% of those on a misread clue, so about 224M could be helped. If we rescue one in five, that is roughly 45M more photos found a month. Estimates, not measurements.
Drop the whole idea if: we ask more than 0.4 questions per search, or fewer than 2 of 3 testers do better than they did on their own.

What it does, and what three people found when they used it

1 YOU TYPE
One sentence, in your own words.

2 IT SHOWS
Every clue it took becomes a chip you can see.

3 YOU FIX IT
Change or drop any chip it read wrong.

4 IT ASKS ONCE
Only when a clue could mean two things.

You type: the bill from last august
It shows you: bill · august, year not given. Each one is a chip you can change or remove.
It asks: did you mean August 2025 or August 2026? Google Photos picks one silently and drops the other.
You get back: the photos it is most sure of. When nothing matches, the closest few and one question.

What we said we would run: three interviewees on their own task, Google Photos first, then here, scored on tries.
What actually ran: three unmoderated sessions here. The tries comparison was not run, so it is untested.
What all three confirmed: the app showing its reading before it searches, and asking rather than guessing. Tester 3 knew the finding before we described it: “That has happened to me.”

VERSION 1 what three people hit
A typed search was silently discarded
The chip still read “year unknown”
Cricket and a market among the bills
36 results for “my sister at the wedding”
A question after the answer was on screen
“Fog on the morning drive”, filed under night

VERSION 2 what is running now
The box keeps what you type
It turns solid and reads “last august”
A date alone can no longer surface a photo
8 results, both wedding photos first
Silent when one result is clearly ahead
The title wins; nine contradictions gone

Every measure: formula, target, and the data it needs

MEASURE	HOW IT IS WORKED OUT	TARGET	DATA WE WOULD NEED
“The one” number	photos found ÷ hunts for something over a year old	**+15pts**	hunt started, photo opened, how old it was
People who type something	searches with a clue in them ÷ all hunts	**48→65%**	what was typed, how many clues
Clues read correctly	clues kept loose ÷ clues that were vague	**95%+**	how each clue was read, and whether it filtered
Photo spotted in results	photo opened ÷ searches that returned something	**70%+**	what was shown, in what order, what was opened
Dead ends rescued	hunts saved ÷ hunts that returned nothing	**30%+**	empty result, question offered, question answered
Questions per search	questions asked ÷ hunts (a limit, not a goal)	**0.4 or less**	question shown, answered, skipped

How to read this: the first row is the one that matters on its own. “Clues read correctly” is what our fix moves directly, and the only one you cannot game by returning more photos. **None of this is instrumented yet —** targets are set against the MVP test, not Google's real numbers, which we do not have.

What could break this, and what we do not yet know

A confident wrong answer hurts more than an empty one
THE WARNING SIGN
People searching again within a day.
WHAT WE WOULD DO, AND ITS COST
Show why each photo matched, and keep the old results alongside. Costs screen space and a slower first load.

The questions turn into an interrogation
THE WARNING SIGN
More than 0.4 questions per search, or half of them skipped.
WHAT WE WOULD DO, AND ITS COST
Only ask when we are genuinely unsure, once per search. When we skip, we show a wider set instead.

Reading your words needs your history
THE WARNING SIGN
People opting out, and privacy complaints.
WHAT WE WOULD DO, AND ITS COST
Run on what is already on the phone, so nothing new leaves it. Costs accuracy if you change devices.

What this research cannot tell you
31 survey replies from people who volunteered. Five interviews. The discovery engine reads public posts, so it hears the people annoyed enough to write one. A 180-item demo is not 75,000 photos, and things are easier to find in a small library. The sizing is a rough order of magnitude.

Honestly: the strongest thing here is one controlled test on one library, and the biggest number comes off a 31-person survey. Everything is sized so it can be proved wrong, rather than sized to impress.